

Marwan Ghazal

Sheikh Zayed, Egypt | +20 101 651 1889 | marwanghazal002y@gmail.com | [LinkedIn](#) | [GitHub](#)

Dear Vionex Team,

Thank you for sponsoring DeepX 2026 and for the paid internship awarded as part of our 3rd-place prize. I'd love to take it up, and I'm writing to formally introduce myself and what I'd bring beyond the hackathon.

At DeepX, my team had 12 hours to build a full Arabic Aspect-Based Sentiment Analysis pipeline: preprocessing, aspect extraction, sentiment classification, and a working end-to-end solution. Arabic NLP is unforgiving (rich morphology, limited data, ABSA on top), and shipping something that actually worked under that pressure was easily the most rewarding part. That small-team, tight-constraint, ship-the-thing mode is what drew me to Vionex: a MENA-focused team building Arabic-language and applied AI that actually gets deployed, not just demoed.

Outside the hackathon, I've been building production AI all year. At the National Training Academy this summer I shipped two systems end-to-end: a CV-to-job ranking service powered by Llama 3.1 + LoRA returning structured JSON for downstream automation, and an Egyptian National ID OCR pipeline (FastAPI + OpenCV + Tesseract) with preprocessing, validation, and a real-time API. Both were deployed to replace manual review in digital-transformation workflows. My research project, Smart Food Guardian, took 1st place at Nile University's Undergraduate Research Forum: a multimodal Nutri-Score classifier built under a strict real-world constraint, using only information physically readable from product packaging, with no leakage from external metadata. Across these projects, I've worked extensively with PyTorch, and recently completed DataCamp's full Deep Learning in Python specialization, covering CNNs, RNNs/LSTMs, and Transformers on real datasets.

One thread that runs through my work is OCR and structured-output pipelines: turning messy real-world inputs (ID cards, product packaging, EEG recordings) into clean, validated JSON that downstream systems can use. My day-to-day stack is Python, FastAPI/Flask, REST/JSON, and the OpenAI/Gemini/Claude APIs, with prompt engineering, structured outputs, embeddings, and RAG as recent extensions. I work quickly in Cursor, Windsurf, Claude code, and Copilot. I'm a junior at Nile University with a 4.0 GPA (President's Honor, Spring 2025), graduating in 2027.

I'd be glad to contribute on LLM integrations, automation pipelines, OCR/document workflows, or backend AI features, wherever the team needs hands. Happy to walk through any of the projects above or take a short technical task to get started.

Looking forward to hearing about next steps.

Best regards,
Marwan Ghazal